

Overview v0 to v5

Source: starter_v0/artifacts/version_log.csv · Base suite (20 cases)

Final outcome: v5 reaches 20/20 pass with case/routing/argument/multiturn accuracy all at 100%.

Metrics by Version

Version	Case acc	Routing acc	Arg acc	Multiturn acc	Passed cases
v0	65%	75%	65%	100%	13/20
v1	80%	85%	80%	100%	16/20
v2	85%	90%	85%	100%	17/20
v3	80%	95%	80%	100%	16/20
v4	90%	90%	90%	83.33%	18/20
v5	100%	100%	100%	100%	20/20

Per-Version Transition Analysis

Transition	Overall takeaway	Detailed changes vs previous	Why better/worse than previous
v0 -> v1	Large quality jump; fewer wrong guesses.	Updated system prompt to enforce clarify-first when required inputs are missing.	v1 is better because it reduces missing-info mistakes and unnecessary direct actions.
v1 -> v2	Steady improvement with stronger boundaries.	Refined prompt with explicit side-effect confirmation and clearer multi-tool execution rule.	v2 is better because routing and argument accuracy both improved (+5 pts each).
v2 -> v3	Regression in overall pass rate despite routing gain.	Shifted emphasis into tool descriptions (tools.yaml) for boundary handling.	v3 is worse overall because argument quality dropped and wrong-tool errors reappeared.
v3 -> v4	Strong recovery in case/argument metrics.	Aligned prompt+tools defaults for missing-info, side-effects, and news/social argument conventions.	v4 is better because it fixes earlier regressions and increases pass count to 18/20.
v4 -> v5	Final step to full pass.	Added two precise prompt constraints: URL-present => fetch directly, and latest-turn narrowing overrides earlier intent.	v5 is better because it removes the last 2 fails and reaches 20/20.